

# Moving Averages in R

William Chiu

5/7/2022

## Data and Noise

Suppose  $y_t$  is a linear function of  $x_t$  and random error:

$$y = \beta_0 + \beta_1 x_t + \epsilon$$

```
set.seed(1)

nobs <- 1000

x <- log(1:nobs) + rnorm(nobs, 0, 0.25)
y <- 2 + 3 * x + rnorm(nobs)

df <- data.frame(t=1:nobs, y=y, x=x)

rm(x, y)
```

But also suppose that  $x_t$  is not directly observable, instead we observe its noisier cousin that contains random errors. For example, the instrument that measures  $x_t$  suffers from random imprecision.

```
df$x_noisy <- df$x + rnorm(nobs,0)
```

Since  $x_t$  is observed with noisy imprecision, this degrades the OLS fit.

```
true_model <- lm(y ~ x, data=df)

noisy_model <- lm(y ~ x_noisy, data=df)
```


	<i>Dependent variable:</i>	
	y	
	(1)	(2)
Constant	1.935 (0.195)***	10.910 (0.304)***
x	3.008 (0.033)***	
x_noisy		1.486 (0.050)***
Observations	1,000	1,000
R <sup>2</sup>	0.895	0.471
Adjusted R <sup>2</sup>	0.895	0.471
Residual Std. Error (df = 998)	1.040	2.337
<i>Note:</i> *p<0.1; **p<0.05; ***p<0.01		

## Visualization

Plot  $x_t$  and  $x_t^{noisy}$  over time.

```
df_long <- pivot_longer(df, -t) %>% filter(name %in% c('x', 'x_noisy'))

ggplot(df_long, aes(x=t, y=value, group=name, color=name, alpha=name)) +
  geom_line(aes(linetype=name)) +
  scale_linetype_manual(values=c("solid", "dashed"))+
  scale_color_manual(values=c('black', 'green')) +
  scale_alpha_manual(values=c(1, 0.25)) +
  theme_bw()
```



## Filter and Moving Averages

We could apply a filter on  $x_t^{noisy}$  to remove some of the “jumpiness”. First, we apply a backward-looking moving average with a  $k$ -period window:

```
k_param <- 12

df$x_noisy_ma <- stats::filter(x=ts(df$x_noisy),
                              filter=rep(1/k_param, k_param),
                              method="convolution",
                              sides=1)

knitr::kable(head(df, k_param+5))
```

t	y	x	x_noisy	x_noisy_ma
1	2.665125	-0.1566135	-1.0427630	NA
2	5.329106	0.7390580	-1.1831969	NA
3	3.798338	0.8897051	2.5094059	NA
4	7.566075	1.7851146	2.3043845	NA
5	7.144840	1.6918149	1.6359649	NA
6	5.097278	1.5866424	2.2830600	NA
7	9.014142	2.0677674	2.1212831	NA
8	6.879722	2.2640227	0.9537392	NA

t	y	x	x_noisy	x_noisy_ma
9	7.776756	2.3411699	0.2181039	NA
10	9.676868	2.2262380	2.0181594	NA
11	9.786649	2.7758406	2.4630540	NA
12	9.530727	2.5823675	1.5241317	1.317111
13	7.606980	2.4096392	2.8268628	1.639579
14	6.805183	2.0853824	1.7699308	1.885673
15	11.318758	2.9892829	3.8148321	1.994459
16	10.109519	2.7613553	4.0526274	2.140146
17	9.896069	2.8291658	2.2040608	2.187487

Second, we apply a “centered” moving average that includes both past and future periods.

```
df$x_noisy_ma_ctr <- stats::filter(x=ts(df$x_noisy),
                                   filter=rep(1/k_param, k_param),
                                   method="convolution",
                                   sides=2)

knitr::kable(head(df, k_param+5))
```

t	y	x	x_noisy	x_noisy_ma	x_noisy_ma_ctr
1	2.665125	-0.1566135	-1.0427630	NA	NA
2	5.329106	0.7390580	-1.1831969	NA	NA
3	3.798338	0.8897051	2.5094059	NA	NA
4	7.566075	1.7851146	2.3043845	NA	NA
5	7.144840	1.6918149	1.6359649	NA	NA
6	5.097278	1.5866424	2.2830600	NA	1.317111
7	9.014142	2.0677674	2.1212831	NA	1.639579
8	6.879722	2.2640227	0.9537392	NA	1.885673
9	7.776756	2.3411699	0.2181039	NA	1.994459
10	9.676868	2.2262380	2.0181594	NA	2.140146
11	9.786649	2.7758406	2.4630540	NA	2.187487
12	9.530727	2.5823675	1.5241317	1.317111	2.184831
13	7.606980	2.4096392	2.8268628	1.639579	2.282482
14	6.805183	2.0853824	1.7699308	1.885673	2.583787
15	11.318758	2.9892829	3.8148321	1.994459	2.693905
16	10.109519	2.7613553	4.0526274	2.140146	2.729823
17	9.896069	2.8291658	2.2040608	2.187487	2.727767

The filters created NA values that should be removed before refitting the models.

```
df_filtered <- df %>% drop_na()

knitr::kable(head(df_filtered))
```

t	y	x	x_noisy	x_noisy_ma	x_noisy_ma_ctr
12	9.530727	2.582368	1.524132	1.317111	2.184831
13	7.606980	2.409639	2.826863	1.639579	2.282482


t	y	x	x_noisy	x_noisy_ma	x_noisy_ma_ctr
14	6.805183	2.085382	1.769931	1.885673	2.583787
15	11.318758	2.989283	3.814832	1.994459	2.693905
16	10.109519	2.761355	4.052627	2.140146	2.729823
17	9.896069	2.829166	2.204061	2.187487	2.727767

## More Visualization

Plot  $x_t$ ,  $x_t^{noisy}$ , and the filtered predictors over time.

```
df_long2 <- pivot_longer(df_filtered, -t) %>%
  filter(name %in% c('x', 'x_noisy', 'x_noisy_ma', 'x_noisy_ma_ctr'))

ggplot(df_long2, aes(x=t, y=value, group=name, color=name, alpha=name)) +
  geom_line(aes(linetype=name)) +
  scale_linetype_manual(values=c("solid", "dashed", "solid", "solid")) +
  scale_color_manual(values=c('black', 'green', 'purple', 'blue')) +
  scale_alpha_manual(values = c(0.25, 0.25, 0.75, 1)) +
  theme_bw()
```



## Refit models with MA predictors

```
ma_model <- lm(y ~ x_noisy_ma, data=df_filtered)
ma_ctr_model <- lm(y ~ x_noisy_ma_ctr, data=df_filtered)
```

	Dependent variable:			
	y			
	(1)	(2)	(3)	(4)
Constant	1.935 (0.195)***	10.910 (0.304)***	4.930 (0.285)***	4.229 (0.312)***
x	3.008 (0.033)***			
x_noisy		1.486 (0.050)***		
x_noisy_ma			2.508 (0.047)***	
x_noisy_ma_ctr				2.613 (0.052)***
Observations	1,000	1,000	983	983
R <sup>2</sup>	0.895	0.471	0.741	0.724
Adjusted R <sup>2</sup>	0.895	0.471	0.741	0.724
Residual Std. Error	1.040 (df = 998)	2.337 (df = 998)	1.484 (df = 981)	1.533 (df = 981)

Note:

\*p<0.1; \*\*p<0.05; \*\*\*p<0.01

## Optimal Filter via 5-fold CV MSE

The hyper-parameter `k_param` controls the smoothness of the `ma` predictor. We could choose the optimal `k_param` by trying different values and measuring the 5-fold cross-validation error for each value. The code below finds the optimal `k_param` for the backward-looking moving average.

```
cv_5_fold <- function(dataframe, ma_window){
  dataframe[, 'x_noisy_ma'] <- stats::filter(x=ts(dataframe[, 'x_noisy']),
                                             filter=rep(1/ma_window, ma_window),
                                             method="convolution",
                                             sides=1)

  train <- dataframe %>% drop_na()

  mod <- glm(y ~ x_noisy_ma, data=train)

  set.seed(123)

  cv_mod <- boot::cv.glm(train, mod, K=5)

  return(data.frame(ma_window=ma_window, MSE=cv_mod$delta[1]))
}
```


Let's try the moving average windows from 2 to 120 (in increments of 2):

```
tuning_ma_window <- lapply(seq(2,120,2), function(x){
  cv_5_fold(dataframe=df, ma_window=x)
})

tuning_ma_window_df <- bind_rows(tuning_ma_window)
```

Plot the window against 5-fold CV MSE:

```
ggplot(tuning_ma_window_df, aes(x=ma_window,y=MSE)) +
  geom_line() +
  ylab("5-fold CV MSE") +
  theme_bw()
```



There does not appear to be a meaningful change in CV MSE after  $k=24$ . Hence, we choose 24 as the optimal moving average window.


```
k_param <- 24

df$x_noisy_ma_cv <- stats::filter(x=ts(df$x_noisy),
  filter=rep(1/k_param, k_param),
  method="convolution",
  sides=1)

df_filtered_cv <- df %>% drop_na()
```

```
df_long_cv <- pivot_longer(df_filtered_cv, -t) %>%
  filter(name %in% c('x', 'x_noisy', 'x_noisy_ma', 'x_noisy_ma_cv'))

ggplot(df_long_cv, aes(x=t, y=value, group=name, color=name, alpha=name)) +
  geom_line(aes(linetype=name)) +
  scale_linetype_manual(values=c("solid", "dashed", "solid", "solid"))+
  scale_color_manual(values=c('black', 'green', 'purple', 'blue')) +
  scale_alpha_manual(values=c(0.25, 0.25, 0.75, 1)) +
  theme_bw()
```



```
ma_model_cv <- lm(y ~ x_noisy_ma_cv, data=df_filtered_cv)
```



Dependent variable:				
	y			
	(1)	(2)	(3)	(4)
Constant	1.935 (0.195)***	10.910 (0.304)***	4.930 (0.285)***	4.709 (0.282)***
x	3.008 (0.033)***			
x_noisy		1.486 (0.050)***		
x_noisy_ma			2.508 (0.047)***	
x_noisy_ma_cv				2.557 (0.047)***
Observations	1,000	1,000	983	971
R <sup>2</sup>	0.895	0.471	0.741	0.756
Adjusted R <sup>2</sup>	0.895	0.471	0.741	0.756
Residual Std. Error	1.040 (df = 998)	2.337 (df = 998)	1.484 (df = 981)	1.350 (df = 969)

Note:

\*p<0.1; \*\*p<0.05; \*\*\*p<0.01

## Filter and Exponential Smoothing

Moving averages remove jumpiness in a time series. They are easy to calculate and understand. However, they create missing values in the beginning and/or end of the time series.

Simple exponential smoothing (SES) also removes jumpiness but does not create any missing values.

Hyndman explains SES [here](#) and [here](#). The smoothing equation is customized to our example:

$$x_t^{ses} = \alpha x_t^{noisy} + (1 - \alpha)x_{t-1}^{ses}$$

If  $t=1$ , the value of  $x_0^{ses}$  is not obvious. Fortunately, the `ets` function estimates both  $x_0^{ses}$  and  $\alpha$  using maximum likelihood.

```
x_ets <- ets(df$x_noisy, "ANN")
coef(x_ets)
```

```
##      alpha      1
## 0.05870055 2.19034239
```

The estimate of the smoothing parameter ( $\alpha$ ) is 0.0587005 and the estimate of  $x_0^{ses}$  is 2.1903424.

Notice that the smoothed values do not contain any NA.

```
df$x_noisy_ses <- x_ets$states[-1,1]
knitr::kable(head(df) %>% select(-x_noisy_ma_ctr, -x_noisy_ma))
```

t	y	x	x_noisy	x_noisy_ma_cv	x_noisy_ses
1	2.665125	-0.1566135	-1.042763	NA	2.000557
2	5.329106	0.7390580	-1.183197	NA	1.813669


t	y	x	x_noisy	x_noisy_ma_cv	x_noisy_ses
3	3.798338	0.8897051	2.509406	NA	1.854509
4	7.566075	1.7851146	2.304384	NA	1.880917
5	7.144840	1.6918149	1.635965	NA	1.866538
6	5.097278	1.5866424	2.283060	NA	1.890988

## More Visualization

```
df_long3 <- pivot_longer(df, -t) %>%
  filter(name %in% c('x', 'x_noisy', 'x_noisy_ma_cv', 'x_noisy_ses'))

ggplot(df_long3, aes(x=t, y=value, group=name, color=name, alpha=name)) +
  geom_line(aes(linetype=name)) +
  scale_linetype_manual(values=c("solid", "dashed", "solid", "solid")) +
  scale_color_manual(values=c('black', 'green', 'purple', 'blue')) +
  scale_alpha_manual(values=c(0.25, 0.25, 0.75, 1)) +
  theme_bw()
```

## Warning: Removed 23 row(s) containing missing values (geom\_path).



## Refit Models

```
ma_model_ets <- lm(y ~ x_noisy_ses, data=df)
```

<i>Dependent variable:</i>				
	y			
	(1)	(2)	(3)	(4)
Constant	1.935 (0.195)***	10.910 (0.304)***	4.709 (0.282)***	4.498 (0.234)***
x	3.008 (0.033)***			
x_noisy		1.486 (0.050)***		
x_noisy_ma_cv			2.557 (0.047)***	
x_noisy_ses				2.600 (0.039)***
Observations	1,000	1,000	971	1,000
R <sup>2</sup>	0.895	0.471	0.756	0.814
Adjusted R <sup>2</sup>	0.895	0.471	0.756	0.814
Residual Std. Error	1.040 (df = 998)	2.337 (df = 998)	1.350 (df = 969)	1.386 (df = 998)

*Note:*

\*p<0.1; \*\*p<0.05; \*\*\*p<0.01

## Conclusion

Filtering a noisy predictor variable ( $x_t^{noisy}$ ) may improve your OLS fit.